

# Project 2: Gridworld Policy Optimization with Dynamic Programming and Monte Carlo Algorithms

**Student:** Harshani Rathnayake  
**Departments:** Mathematics and Statistics and Computer Science  
**Date:** July 21, 2025  
**GitHub Repository:** <https://github.com/Harshani-Rathnayake/Project-2-Gridworld-Policy-Optimization-with-Dynamic-Programming-and-Monte-Carlo-Algorithms.git>

---

## Abstract

This study investigates reinforcement learning methods in a basic  $5 \times 5$  gridworld setting. First, it analyzes the Bellman equations directly and evaluates the policy iteratively in order to estimate state value functions under a uniform random policy. Next, this study uses explicit solutions of the Bellman optimality formula, policy iteration, and value iteration to identify optimal policies. Moreover, it investigates Monte Carlo methods with and without investigating begins, and off-policy learning with importance sampling in an enlarged version of the environment containing terminal states. The results show which states have the highest expected returns and show how various algorithms converge to reliable answers.

## Introduction

Training an agent to interact with an environment and make successive decisions that maximize cumulative rewards is a challenge that reinforcement learning attempts to solve. The gridworld environment is a common standard for researching and contrasting reinforcement learning algorithms because of its ease of use and interpretability. In this study, a  $5 \times 5$  gridworld with special states, distinct reward dynamics, and transitions is analyzed. Using traditional dynamic programming techniques, this study determines optimal policies and estimates value functions under a uniform random policy. Furthermore, it uses off-policy learning with importance sampling and Monte Carlo methods to study learning in episodic settings with terminal states. This project shows and contrasts the convergence qualities and efficacy of different methods in identifying the best ways to make decisions.

## Environment and Problem Setup

The  $5 \times 5$  grid that makes up the environment produces 25 distinct states. The agent has four options for each state: move up, down, left, or right. Movements are controlled by the grid borders; attempting to step off the grid results in no change in position and a negative reward.

## Part 1: Original gridworld

Initially, the grid contains four distinctive colored squares:

- **Blue square** - first row, second column: Any action results in an immediate jump to the red square and yields a reward of +5.
- **Green square** - first row, fifth column: Any action leads to a stochastic transition to either the red or yellow square, each with probability 0.5, yielding a reward of +2.5.
- **Red square** - fourth row, third column: A target state reached from the blue and green squares.
- **Yellow square** - fifth row, fifth column: Another target state is reachable from the green square.

All moves from white, red, or yellow squares earn 0; however, attempting to step outside the grid from a white or yellow square results in a penalty of  $-0.5$ . Future reward present values are calculated using a discount factor of  $\gamma = 0.95$ .




Figure 1: Original gridworld environment showing special colored squares

## Part 2: Modified gridworld with terminal states

In the second part, the gridworld is modified as follows:

- The positions of the special squares are updated:
  - **Blue square** - first row, second column
  - **Green square** - first row, fifth column
  - **Red square** - fifth row, third column
  - **Yellow square** - fifth row, fifth column
- Three black terminal squares are added:
  - third row, first column

- third row, fifth column
- fifth row, first column

The episode terminates instantly when the agent enters any black terminal square.

Any move from a white, yellow, or red square to any other square in this configuration results in a reward of  $-0.2$ , while trying to leave the grid from these squares still results in a penalty of  $-0.5$ . The same  $\gamma = 0.95$  discount factor is applied.




Figure 2: Modified gridworld environment showing special colored squares and black terminal squares

## Part 1: Policy Evaluation and Optimal Policy Computation

This section initially examines the behavior of an agent in accordance with a uniform random policy, in which the four possible actions (up, down, left, and right) are all equally likely, in order to create a performance baseline. Under this policy, estimating the state value function reveals which states, even in the absence of learning, naturally produce higher long-term returns. This baseline is used as a point of comparison and evaluation for the efficacy of later, more complex policy optimization techniques.

### Estimating the Value Function under a Uniform Random Policy

According to the uniform random policy, every action in every state has the same probability (0.25). The state value function was estimated using two methods under this policy:

**Explicit Solution of Bellman Equations:** The probabilities of policy action were used to weight transitions and rewards in order to create the expected transition probability matrix  $P_\pi$  and the expected reward vector  $R_\pi$ . The vector of the value function  $\mathbf{V}$  was computed by solving the linear system:

$$\mathbf{V} = (I - \gamma P_\pi)^{-1} R_\pi,$$

where  $I$  is the identity matrix.

**Iterative Policy Evaluation:** The Bellman expectation backup was applied iteratively until convergence, updating each state value in the process described below, beginning with an initial value function of zeros:

$$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V(s')].$$

The process ended when the maximum change in state values was below a threshold ( $10^{-6}$ ).

**Results:** The two approaches showed consistency as they converged to the same value estimates. The estimate of the state value function is shown in Table 1.

2.17	<b>4.73</b>	2.07	1.27	1.78
1.12	1.78	1.17	0.74	0.56
0.16	0.48	0.35	0.11	-0.19
-0.55	-0.28	-0.28	-0.44	-0.74
-1.11	-0.85	-0.81	-0.94	-1.24

Table 1: Estimated state value function under uniform random policy ( $\gamma = 0.95$ )

The blue square has the highest value (4.73), which is consistent with its deterministic jump to the red square and high immediate reward.

## Deriving the Optimal Policy

Three methods were used to create optimal policies and value functions: value iteration, policy iteration, and explicit solution of the Bellman optimality equation mentioned above, all with the discount factor  $\gamma = 0.95$ .

**Policy Iteration:** In order to reach convergence, policy iteration alternates between policy evaluation and greedy policy improvement. The resulting value function and policy are shown in Tables 2 and 3.

**Value Iteration:** The value iteration selects the optimum action for each state and then iteratively applies the Bellman optimality update until convergence, at which point the optimal policy is retrieved.

**Results:** Consistent optimal value functions and policies were generated by both approaches, as demonstrated below.

21.00	22.10	21.00	19.95	18.38
19.95	21.00	19.95	18.95	18.00
18.95	19.95	18.95	18.00	17.10
18.00	18.95	18.00	17.10	16.25
17.10	18.00	17.10	16.25	15.43

Table 2: Optimal state value function ( $\gamma = 0.95$ ) from policy and value iteration

3	0	2	2	0
0	0	0	0	2
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Table 3: Optimal policy actions (0: Up, 1: Down, 2: Left, 3: Right)

**Policy Visualization:** Figures 3 and 4 illustrate the derived optimal policies from policy iteration and value iteration, respectively.




Figure 3: Optimal policy derived via policy iteration. Arrows represent recommended actions.

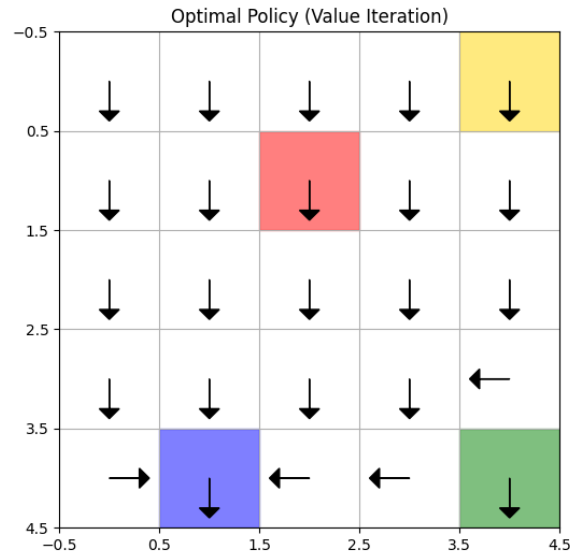


Figure 4: Optimal policy derived via value iteration. Arrows represent recommended actions.

Both methods converge to the same policy, confirming consistency under the defined environment.

## Part 1: Discussion and Analysis

The outcomes of the policy optimization and evaluation process provide important new information on the structure of the gridworld environment and the efficacy of the reinforcement learning techniques used.

First, the blue square has the highest predicted return (4.73), according to the estimated state value function under the uniform random policy. This is consistent with intuition: the red square deterministically replaces the blue square, giving an instant reward of +5, which is then gradually discounted and spread to nearby states. In comparison, the green square has a lower expected return despite being a special state because of its stochastic nature: transitions divide between the red and yellow squares, effectively diluting the immediate reward to an average of +2.5.

The accuracy of the implementation and the coherence of the mathematical framework were demonstrated by the same value estimations obtained from the explicit solution of the Bellman equations and the iterative policy evaluation. The minor negative values in the lower rows represent states that, on average, are more likely to result in off-grid moves and penalties before reaching higher value states, resulting in slightly negative expected returns.

In the second part of the study, the optimal policies produced by policy iteration and value iteration converged to the same result, validating the theoretical hypothesis that both methods should provide the same optimal policy when the environment and discount factor remain fixed. The optimal value function yields values much higher than the uniform random policy, with a peak value of 22.10 in the blue square. This increase shows that the agent has the ability to exploit high-reward transitions repeatedly while avoiding penalties by adhering to the optimal policy.[1]

Moving toward the blue square and then using the deterministic jump to the red square is the main way that the optimal policy itself efficiently directs the agent towards the rewarding states. By prioritizing behaviors that result in immediate and frequent high rewards, the agent learns to maximize cumulative reward. This pattern demonstrates the usefulness of reinforcement learning algorithms in finding efficient strategies even in basic contexts.[1]

All things considered, the convergence of several approaches to the same optimal policy and value function demonstrates the adaptability of the reinforcement learning technique used. These findings also show that agent behavior is significantly influenced by incentive and environment design, resulting in policies that prioritize approaches with higher expected or immediate returns.

## Part 2: Monte Carlo Methods in Episodic Gridworld

To estimate state value functions and learn policies, Monte Carlo methods were used to sample episodes in the same  $5 \times 5$  gridworld environment, which is now considered episodic. After a certain number of steps or when the agent reaches a terminal state (special colored squares), an episode comes to an end.

### On-Policy Monte Carlo Methods

Two on-policy methods have been examined below:

**Exploring Starts:** Each episode starts from a random state-action pair using the Monte Carlo with exploring starts method in order to promote varied exploration. The estimated value function and the learned policy were as follows after numerous episodes:

$$\begin{bmatrix} 2 & 0 & 1 & 0 & 0 \\ 1 & 1 & 1 & 3 & 1 \\ 0 & 2 & 3 & 3 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 0 & 2 & 1 & 1 & 0 \end{bmatrix}$$

Monte Carlo with exploring starts method - learned policy (0:Up, 1:Down, 2:Left, 3:Right)

$$\begin{bmatrix} -0.5 & -4.21 & -0.73 & -0.62 & -1.38 \\ -0.2 & -0.37 & -0.59 & -0.49 & -0.2 \\ 0.0 & -0.2 & -0.4 & -0.2 & 0.0 \\ -0.2 & -0.39 & -0.59 & -0.38 & -0.2 \\ 0.0 & -0.2 & -0.5 & -0.5 & -0.39 \end{bmatrix}$$

Value function from Monte Carlo with exploring starts method

**Monte Carlo with  $\varepsilon$ -soft policy:** This strategy encourages more extensive exploration by having the agent adopt a  $\varepsilon$ -soft policy, which is mostly greedy with occasional random actions. The learned policy and value function generated are shown below:

$$\begin{bmatrix} 3 & 1 & 2 & 2 & 1 \\ 0 & 3 & 0 & 2 & 2 \\ 0 & 0 & 1 & 0 & 0 \\ 3 & 0 & 0 & 0 & 2 \\ 0 & 0 & 1 & 0 & 2 \end{bmatrix}$$

Monte Carlo with  $\varepsilon$ -soft policy method - learned policy (0:Up, 1:Down, 2:Left, 3:Right)

$$\begin{bmatrix} 3.05 & 2.83 & 2.34 & 1.77 & 1.72 \\ 2.49 & 1.71 & 1.53 & 1.20 & 1.32 \\ 0.00 & 1.17 & 1.01 & 0.70 & 0.00 \\ 0.81 & 0.50 & 0.20 & 0.15 & 0.30 \\ 0.00 & 0.13 & -0.77 & -0.38 & 0.15 \end{bmatrix}$$

Value function from Monte Carlo with  $\varepsilon$ -soft method

## Off-Policy Learning Using Importance Sampling

An off-policy Monte Carlo method was used to evaluate a deterministic target policy while following a different behavior policy, applying importance sampling to correct for the distribution mismatch. The resulting learned policy was:

$$\begin{bmatrix} 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

## Visualization and Comparison

Figure 5 displays the policies learned by each Monte Carlo method, with arrows representing the chosen actions and special squares marked with colors, while black squares represent terminal states.




Figure 5: Learned policies from Monte Carlo methods: (Left) Exploring Starts; (Center)  $\epsilon$ -soft; (Right) Off-policy with importance sampling.

Table 4 summarizes the estimated value functions from each method, indicating differences in learning quality and convergence behavior.

Method	Characteristics of Value Estimates
Exploring Starts	Due to high variance and sparse reward input, the values are primarily negative with scattered low values, suggesting unstable learning and poor convergence.
$\epsilon$ -soft Policy	Higher, smoother values indicate better exploration-exploitation balance and learning stability, particularly within close proximity of rewarding situations.
Off-Policy with Importance Sampling	Sparse and less reliable estimates; large variance in importance weights leads to unstable value estimation and fragmented policies.

Table 4: Summary of value function estimates from Monte Carlo methods

## Part 2: Discussion and Interpretation

The Monte Carlo experiments highlight several important aspects of learning in episodic gridworld environments:

- **Exploring Starts** promotes extensive exploration by beginning each episode at random, but it has poor sample usage and large return variation, which leads to erratic and scattered policy and value estimations.
- **$\epsilon$ -soft policies** enhance learning by the combination of stochastic action selection and mostly greedy behavior. As a result, the value function converges more consistently, and state space coverage is improved. The agent is typically effectively guided toward high-reward states by the learned policies.



- **Off-policy Monte Carlo method**, Although theoretically strong, these face real-world difficulties in this setting because there are not many ideal action scenes. As a result, importance sampling ratios increase, resulting in unstable updates and large variance.

The observed variations among approaches highlight the need to establish a balance between exploration and exploitation as well as the challenges of off-policy evaluation using importance sampling, particularly in settings with lengthy episodes and few rewards.[2]

These results indicate that on-policy Monte Carlo methods with controlled exploration may yield more consistent and useful learning outcomes for episodic tasks with limited rewards. In contrast, off-policy methods might require the use of different algorithms or additional variance reduction strategies.

## Conclusion

This study focused on both policy evaluation and optimal policy determination, demonstrating the application of multiple reinforcement learning techniques to a  $5 \times 5$  gridworld setting. Iterative policy evaluation, policy iteration, value iteration, and explicit Bellman equation solving are examples of dynamic programming techniques that consistently generate stable and convergent value functions and policies by using the dynamics of the given environment.

Monte Carlo methods used in the episodic scenario provided useful insights into learning from the sampling experience. The significance of establishing a balance between exploration and exploitation is shown by the reasonable policy quality and smoother value estimations obtained by on-policy Monte Carlo methods with  $\epsilon$ -soft exploration. On the other hand, off-policy Monte Carlo with importance sampling had challenges because of variance explosion, which limited its practical usefulness in this setting, whereas exploratory Monte Carlo experienced large variance and unstable learning.

Overall, the findings highlight the advantages and disadvantages of model-based versus model-free approaches: Monte Carlo methods provide flexible learning from experience but require careful exploration strategies and variance management, while dynamic programming performs best when fully aware of transition dynamics. Even in straightforward gridworld settings, this analysis highlights how important environment structure, reward design, and algorithmic selection are for reinforcement learning performance and stability.

# Appendix

## Part 1: Dynamic Programming in Gridworld

```
1
2 #Defined gridworld environment:
3 grid_size = 5 #5x5 grid
4 gamma = 0.95 #reward discount factor
5 actions = ['U', 'D', 'L', 'R'] #action space: Up, Down, Left, Right
6 action_prob = 0.25 #uniform random policy: equal probability for each action
7
8 #Helper: This convert (row, col) to state index
9 def to_state(row, col):
10     return row * grid_size + col
11
12 #Special states defined by the problem statement
13 blue = (0,1) #first row, second column
14 green = (0,4) #first row, fifth column
15 red = (3,2) #fourth row, third column
16 yellow = (4,4) #fifth row, fifth column
17
18 #Converts the special states to indices:
19 blue_s = to_state(*blue)
20 green_s = to_state(*green)
21 red_s = to_state(*red)
22 yellow_s = to_state(*yellow)
23
24 #Total number of states
25 n_states = grid_size * grid_size
26
27 #Initialized transition probability (P) and reward (R) tensors
28 #Dimensions: [state, action, next_state]
29 P = np.zeros((n_states, len(actions), n_states))
30 R = np.zeros((n_states, len(actions), n_states))
31
32 #Build environment dynamics (transition + reward model)
33 for row in range(grid_size):
34     for col in range(grid_size):
35         s = to_state(row, col)
36         for a_idx, a in enumerate(actions):
37             #Special state: blue square always jumps to red with reward 5
38             if s == blue_s:
39                 next_s = red_s
40                 P[s,a_idx,next_s] = 1.0
41                 R[s,a_idx,next_s] = 5
42             #Special state: green square jumps randomly to yellow or red
43             #with reward 2.5
44             elif s == green_s:
45                 P[s,a_idx,red_s] = 0.5
46                 P[s,a_idx,yellow_s] = 0.5
47                 R[s,a_idx,red_s] = 2.5
48                 R[s,a_idx,yellow_s] = 2.5
49             else:
50                 #Regular square: compute next position after action
51                 if a == 'U':
52                     next_row = max(row-1,0)
53                     next_col = col
54                 elif a == 'D':
55                     next_row = min(row+1,grid_size-1)
56                     next_col = col
```

```

56         elif a == 'L':
57             next_row = row
58             next_col = max(col-1,0)
59         elif a == 'R':
60             next_row = row
61             next_col = min(col+1,grid_size-1)
62         next_s = to_state(next_row, next_col)
63         #Below checks if move hits border:
64         if next_row == row and next_col == col and (
65             (row==0 and a=='U') or
66             (row==grid_size-1 and a=='D') or
67             (col==0 and a=='L') or
68             (col==grid_size-1 and a=='R')):
69             P[s,a_idx,next_s] = 1.0
70             R[s,a_idx,next_s] = -0.5 #penalty for trying to step
71         off grid
72         else:
73             P[s,a_idx,next_s] = 1.0
74             R[s,a_idx,next_s] = 0 #normal move with zero reward
75
76     #(1) Solve Bellman equations explicitly
77     #Compute expected transition matrix P_pi and expected reward R_pi under
78     uniform policy
79     P_pi = np.zeros((n_states, n_states))
80     R_pi = np.zeros(n_states)
81     for s in range(n_states):
82         for a_idx in range(len(actions)):
83             for next_s in range(n_states):
84                 P_pi[s,next_s] += action_prob * P[s,a_idx,next_s]
85                 R_pi[s] += action_prob * P[s,a_idx,next_s] * R[s,a_idx,next_s]
86
87     #Solve linear system:  $V = (I - \gamma * P_{\pi})^{-1} * R_{\pi}$ 
88     I = np.eye(n_states)
89     V_explicit = np.linalg.solve(I - gamma * P_pi, R_pi)
90     print("Value function by solving Bellman equations:")
91     print(np.round(V_explicit.reshape((grid_size,grid_size)),2))
92
93     #Find highest value states
94     max_val = np.max(V_explicit)
95     max_states = np.argwhere(np.isclose(V_explicit, max_val)).flatten()
96     positions = [(s // grid_size, s % grid_size) for s in max_states]
97     print(f"\nHighest value: {max_val:.2f}")
98     print("States with highest value (index):", max_states)
99     print("Positions (row,col):", positions)
100
101     #(2) Iterative policy evaluation under uniform random policy
102     #Goal: Estimate value function V(s) assuming the agent picks each action
103     with equal probability
104     V = np.zeros(n_states) #Initialize value function: all zeros
105     theta = 1e-6 #Convergence threshold: stop when changes are very small
106     while True:
107         delta = 0 #Track the maximum change in V across all states
108         V_new = np.zeros_like(V) #Temporary array to store new value estimates
109         #Loop over all states
110         for s in range(n_states):
111             v = 0 #To accumulate expected return for state s under uniform
112             random policy
113             #Loop over all possible actions

```

```

112     for a_idx in range(len(actions)):
113         #Loop over all possible next states
114         for next_s in range(n_states):
115             #Expected contribution:
116             #action_prob      transition probability      [reward +
value of next state]
117             v += action_prob * P[s, a_idx, next_s] * (R[s, a_idx, next_s
] + gamma * V[next_s])
118             V_new[s] = v #Update value for state s
119             #Update delta to track the largest difference from previous
iteration
120             delta = max(delta, abs(V_new[s] - V[s]))
121             V = V_new #Update value function for next iteration
122             #Check for convergence: if max change is below threshold, stop
123             if delta < theta:
124                 break
125 #Display final estimated value function
126 print("\nValue function by iterative policy evaluation:")
127 print(np.round(V.reshape((grid_size, grid_size)), 2))
128
129 #(3) Policy iteration: alternating between policy evaluation and improvement
130 #Initialize policy: start by always taking action 0
131 policy = np.zeros(n_states, dtype=int)
132 #Initialize value function: start with all zeros
133 V_pi = np.zeros(n_states)
134 while True:
135     #Step 1: Policy Evaluation
136     #Compute value function V_pi for the current fixed policy until it
converges
137     while True:
138         delta = 0 #Track maximum change to check convergence
139         V_new = np.zeros_like(V_pi) #Temporary array for updated values
140         #Loop over each state to update value
141         for s in range(n_states):
142             a_idx = policy[s] #action chosen by current policy in state s
143             #Compute expected return:
144             #sum over next states of P(s,a,next_s) * (reward +      * V[next_s
])
145             v = sum(
146                 P[s, a_idx, next_s] * (R[s, a_idx, next_s] + gamma * V_pi[
next_s])
147                 for next_s in range(n_states)
148             )
149             V_new[s] = v
150             #Update delta to track largest change
151             delta = max(delta, abs(V_new[s] - V_pi[s]))
152             V_pi = V_new #update value function
153             #Stop policy evaluation if value function converged (small changes
only)
154             if delta < theta:
155                 break
156
157     #Step 2: Policy Improvement
158     #Update policy to be greedy w.r.t. current value function V_pi
159     policy_stable = True #Flag to check if policy changes in this step
160     for s in range(n_states):
161         old_action = policy[s] #remember old action to compare later
162         #Compute expected return for each possible action
163         action_values = np.zeros(len(actions))
164         for a_idx in range(len(actions)):

```

```

165         action_values[a_idx] = sum(
166             P[s, a_idx, next_s] * (R[s, a_idx, next_s] + gamma * V_pi[
next_s]))
167         for next_s in range(n_states)
168     )
169     #Choose the action with highest expected return
170     best_action = np.argmax(action_values)
171     policy[s] = best_action
172     #If policy changed, mark as unstable so we keep iterating
173     if old_action != best_action:
174         policy_stable = False
175     #If policy didn't change for any state, we have found the optimal policy
176     if policy_stable:
177         break
178 #Display final results
179 print("\nOptimal value function (policy iteration):")
180 print(np.round(V_pi.reshape((grid_size,grid_size)),2))
181 print("\nOptimal policy (0:U,1:D,2:L,3:R):")
182 print(policy.reshape((grid_size,grid_size)))
183
184 #(4) Value iteration to compute optimal value function and policy
185 #Initialize value function: start with all zeros
186 V_vi = np.zeros(n_states)
187 while True:
188     delta = 0      #Tracks max change in value function to check for
convergence
189     V_new = np.zeros_like(V_vi) #Temporary array to store new value
estimates
190     #Loop over each state to update its value
191     for s in range(n_states):
192         action_values = np.zeros(len(actions)) #Store expected return for
each possible action
193         #For each action, compute expected return:
194         #sum over all possible next states: P(s,a,next_s) * [ R(s,a,next_s)
+
* V(next_s) ]
195         for a_idx in range(len(actions)):
196             action_values[a_idx] = sum(
197                 P[s, a_idx, next_s] * (R[s, a_idx, next_s] + gamma * V_vi[
next_s]))
198             for next_s in range(n_states)
199         )
200         #Choose the maximum expected return among actions      Bellman
optimality update
201         V_new[s] = np.max(action_values)
202         #Update delta to see how much V changed
203         delta = max(delta, abs(V_new[s] - V_vi[s]))
204     #Update value function
205     V_vi = V_new
206     #Check for convergence: if changes are smaller than small threshold
theta, stop
207     if delta < theta:
208         break
209
210 #Extract optimal policy from final value function
211 policy_vi = np.zeros(n_states, dtype=int) #Array to store best action per
state
212 for s in range(n_states):
213     action_values = np.zeros(len(actions))
214     #For each action, compute expected return as before
215     for a_idx in range(len(actions)):

```

```

216         action_values[a_idx] = sum(
217             P[s, a_idx, next_s] * (R[s, a_idx, next_s] + gamma * V_vi[next_s
218         ])
219         for next_s in range(n_states)
220     )
221     #Best action is the one with highest expected return
222     policy_vi[s] = np.argmax(action_values)
223 #Display results
224 print("\nOptimal value function (value iteration):")
225 print(np.round(V_vi.reshape((grid_size,grid_size)),2))
226 print("\nOptimal policy (0:U,1:D,2:L,3:R):")
227 print(policy_vi.reshape((grid_size,grid_size)))
228 #Plot function to visualize policy as arrows
229 def plot_policy(policy, title):
230     fig, ax = plt.subplots(figsize=(6,6))
231     ax.set_xlim(-0.5, grid_size-0.5)
232     ax.set_ylim(-0.5, grid_size-0.5)
233     ax.set_xticks(np.arange(-0.5, grid_size, 1))
234     ax.set_yticks(np.arange(-0.5, grid_size, 1))
235     ax.grid(True)
236     ax.set_title(title)
237     #Draw special squares
238     special_states = {
239         'blue': (blue, 'blue'),
240         'green': (green, 'green'),
241         'red': (red, 'red'),
242         'yellow': (yellow, 'gold')
243     }
244     for name, (pos, color) in special_states.items():
245         row, col = pos
246         ax.add_patch(plt.Rectangle((col-0.5, grid_size - row -1 -0.5), 1, 1,
247             color=color, alpha=0.5))
248     #Draw arrows for each state's chosen action
249     action_to_delta = {
250         0: (0, +0.3), #U
251         1: (0, -0.3), #D
252         2: (-0.3, 0), #L
253         3: (+0.3, 0) #R
254     }
255     for row in range(grid_size):
256         for col in range(grid_size):
257             s = to_state(row,col)
258             a = a = policy[s] if np.isscalar(policy[s]) else np.argmax(
259                 policy[s])
260             dx, dy = action_to_delta[a]
261             plot_row = grid_size - row -1
262             ax.arrow(col, plot_row, dx, dy, head_width=0.2, head_length=0.1,
263                 fc='k', ec='k')
264             plt.gca().invert_yaxis()
265             plt.show()
266 #Plots the optimal policies found
267 plot_policy(policy, "Optimal Policy (Policy Iteration)")
268 plot_policy(policy_vi, "Optimal Policy (Value Iteration)")
269 #Compare policies: are they identical or not
270 same = np.all(policy == policy_vi)
271 print("Are policies from policy iteration and value iteration identical?:",
272     same)

```

```

271 #Show where they differ:
272 diff = (policy != policy_vi).reshape((grid_size, grid_size))
273 print("Differences between policies (True = different):")
274 print(diff)

```

Listing 1: Python Code for Value Iteration, Policy Iteration and Bellman Equations in Gridworld

## Part 2: Monte Carlo Methods in Gridworld

```

1
2 #Define gridworld and parameters
3 grid_size = 5 #Gridworld is 5x5 cells
4 gamma = 0.95 #Discount factor: future rewards are worth 95% as much
5 actions = ['U', 'D', 'L', 'R'] #Action set: Up, Down, Left, Right
6 action_idx = {a:i for i,a in enumerate(actions)} #Helper: map actions to
    indices for array indexing
7
8 #Define special colored squares (positions are (row, col) starting from 0)
9 blue = (0, 1) #first row, second column
10 green = (0, 4) #first row, fifth column
11 red = (4, 2) #last row, third column
12 yellow = (4, 4) #last row, last column
13 #Terminal (black) squares where episode ends immediately when agent steps in
14 black_terminals = [
15     (2, 0), #third row, first column
16     (2, 4), #third row, last column
17     (4, 0) #last row, first column
18 ]
19
20 #Convert positions (row, col) to state indices (flattened index)
21 blue_s = blue[0]*grid_size + blue[1]
22 green_s = green[0]*grid_size + green[1]
23 red_s = red[0]*grid_size + red[1]
24 yellow_s = yellow[0]*grid_size + yellow[1]
25 black_s = [r*grid_size + c for (r, c) in black_terminals] #list of terminal
    state indices
26 n_states = grid_size * grid_size #total number of states in grid
27
28 #Initialize starting policy: uniform random policy
29 #Each state: equal probability for each action
30 policy = np.ones((n_states, len(actions))) / len(actions)
31
32 #Function to simulate a single episode in the Gridworld environment under a
    given policy
33 #Parameters:
34 #policy : current policy, mapping each state to a probability distribution
    over actions
35 #exploring_starts : if True, start from a random non-terminal state and
    random action
36 #max_steps : maximum number of steps to run in an episode to prevent
    infinite loops
37 #episode : list of (state, action, reward) tuples representing the
    trajectory
38 def generate_episode(policy, exploring_starts=False, max_steps=100):
39     episode = [] #list to store the sequence of (state, action, reward)
40     #Choose starting state and starting action
41     if exploring_starts:
42         #For exploring starts: start from a random non-terminal state and a
        random action
43         s = random.choice([s for s in range(n_states) if s not in black_s])

```

```

44     a = random.choice(range(len(actions)))
45 else:
46     #Otherwise: start from a random non-terminal state and select action
    according to the current policy
47     s = random.choice([s for s in range(n_states) if s not in black_s])
48     a = np.random.choice(len(actions), p=policy[s])
49 #Loop to simulate the episode, step by step
50 for _ in range(max_steps):
51     #Convert flattened state index to (row, col) coordinates
52     row, col = divmod(s, grid_size)
53     #Check if current state is a terminal state (black square): if yes,
episode ends
54     if s in black_s:
55         break
56     #Check for special colored squares and apply special transitions and
rewards
57     if s == blue_s:
58         reward = 5 #reward for blue square
59         next_s = red_s #teleport to red square
60     elif s == green_s:
61         reward = 2.5 #reward for green square
62         #teleport randomly to red or yellow square
63         next_s = np.random.choice([red_s, yellow_s])
64     else:
65         #For normal white squares: determine next position based on
action
66         if actions[a] == 'U': #move up
67             next_row = max(row-1, 0)
68             next_col = col
69         elif actions[a] == 'D': #move down
70             next_row = min(row+1, grid_size-1)
71             next_col = col
72         elif actions[a] == 'L': #move left
73             next_row = row
74             next_col = max(col-1, 0)
75         else: #move right ('R')
76             next_row = row
77             next_col = min(col+1, grid_size-1)
78         #Check if the agent attempted to step off the grid (hit the wall
)
79         if next_row == row and next_col == col and (
80             (row==0 and actions[a]=='U') or
81             (row==grid_size-1 and actions[a]=='D') or
82             (col==0 and actions[a]=='L') or
83             (col==grid_size-1 and actions[a]=='R')):
84             reward = -0.5 # penalty for hitting the wall
85         else:
86             reward = -0.2 #penalty for moving from white square to white
square
87         #Convert next (row, col) to flattened state index
88         next_s = next_row * grid_size + next_col
89         #Add the current step to the episode list
90         episode.append((s, a, reward))
91         # Move to the next state
92         s = next_s
93         # Check again if the new state is a terminal state: if yes, episode
ends
94         if s in black_s:
95             break
96         #Choose next action in the new state according to the current policy

```



```

97     a = np.random.choice(len(actions), p=policy[s])
98     #Return the full episode: list of (state, action, reward)
99     return episode
100
101 #Helper for policy iteration (not Monte Carlo): build P_pi and R_pi under
    deterministic policy
102 def build_P_pi_and_R_pi(policy, P, R):
103     P_pi = np.zeros((n_states, n_states)) #transition matrix under policy
104     R_pi = np.zeros(n_states) #expected reward vector under policy
105     for s in range(n_states):
106         a = policy[s] #deterministic action chosen in state s
107         P_pi[s, :] = P[s, a, :] #probabilities of next states
108         R_pi[s] = np.sum(P[s, a, :] * R[s, a, :]) #expected immediate
    reward
109     return P_pi, R_pi
110
111 #Policy iteration using explicit matrix inversion (solve Bellman equations)
112 def policy_iteration_explicit(P, R, gamma, policy, max_iters=100):
113     for i in range(max_iters):
114         #Step 1: Policy evaluation
115         P_pi, R_pi = build_P_pi_and_R_pi(policy, P, R)
116         I = np.eye(n_states)
117         V = np.linalg.solve(I - gamma * P_pi, R_pi) #solve  $V = (I - \gamma * P_{\pi})^{-1} * R_{\pi}$ 
118         #Step 2: Policy improvement
119         policy_stable = True
120         for s in range(n_states):
121             action_values = np.zeros(len(actions))
122             for a in range(len(actions)):
123                 #expected return if we take action a in state s
124                 action_values[a] = np.sum(P[s, a, :] * (R[s, a, :] + gamma *
    V))
125             best_action = np.argmax(action_values)
126             if best_action != policy[s]:
127                 policy[s] = best_action
128                 policy_stable = False
129         if policy_stable:
130             print(f"Policy converged after {i+1} iterations")
131             break
132     return policy, V
133
134 #Same plot function as before, but with the black terminal squares
135 def plot_policy(policy, title):
136     fig, ax = plt.subplots(figsize=(6,6))
137     ax.set_xlim(-0.5, grid_size-0.5)
138     ax.set_ylim(-0.5, grid_size-0.5)
139     ax.set_xticks(np.arange(-0.5, grid_size, 1))
140     ax.set_yticks(np.arange(-0.5, grid_size, 1))
141     ax.grid(True)
142     ax.set_title(title)
143     #Draw special squares
144     special_states = {
145         'blue': (blue, 'blue'),
146         'green': (green, 'green'),
147         'red': (red, 'red'),
148         'yellow': (yellow, 'gold')
149     }
150     for name, (pos, color) in special_states.items():
151         row, col = pos
152         ax.add_patch(plt.Rectangle((col-0.5, grid_size - row - 1 - 0.5), 1, 1,

```

```

153     color=color, alpha=0.5))
154 #NEW: Draw black terminal squares
155 for (row, col) in black_terminals:
156     ax.add_patch(plt.Rectangle((col-0.5, grid_size - row -1 -0.5), 1, 1,
157     color='black', alpha=0.7))
158 #Draw arrows for policy
159 action_to_delta = {
160     0: (0, +0.3), #U
161     1: (0, -0.3), #D
162     2: (-0.3, 0), #L
163     3: (+0.3, 0) #R
164 }
165 for row in range(grid_size):
166     for col in range(grid_size):
167         s = row * grid_size + col
168         a = a = policy[s] if np.isscalar(policy[s]) else np.argmax(
169 policy[s])
170         dx, dy = action_to_delta[a]
171         plot_row = grid_size - row -1
172         ax.arrow(col, plot_row, dx, dy, head_width=0.2, head_length=0.1,
173         fc='k', ec='k')
174 plt.gca().invert_yaxis()
175 plt.show()
176
177 # (1) Monte Carlo control with exploring starts (start random state/action)
178 Q = np.zeros((n_states, len(actions))) # Action-value function
179 returns_count = np.zeros((n_states, len(actions))) # Count of visits per (s
180 ,a)
181 for i in range(5000): #number of episodes
182     episode = generate_episode(policy, exploring_starts=True)
183     G = 0 #return (discounted sum of rewards)
184     visited = set() #to only update first visit in each (s,a)
185     for s, a, r in reversed(episode):
186         G = gamma*G + r
187         if (s,a) not in visited:
188             returns_count[s,a] += 1
189             #Incremental mean update
190             Q[s,a] += (G - Q[s,a]) / returns_count[s,a]
191             best_a = np.argmax(Q[s])
192             #Update policy: greedy in Q
193             policy[s] = np.eye(len(actions))[best_a]
194             visited.add((s,a))
195 print("MC with exploring starts - learned policy:")
196 print(np.argmax(policy, axis=1).reshape((grid_size, grid_size)))
197 plot_policy(np.argmax(policy, axis=1), "MC exploring starts")
198 #Compute value function as max_a Q(s,a)
199 V_mc_es = np.max(Q, axis=1)
200 print("\nValue function from MC with exploring starts:")
201 print(np.round(V_mc_es.reshape((grid_size, grid_size)), 2))
202
203 # (2) Monte Carlo control with  $\epsilon$ -soft policy
204 Q2 = np.zeros((n_states, len(actions)))
205 returns2 = np.zeros((n_states, len(actions)))
206 epsilon = 0.1
207 policy2 = np.ones((n_states, len(actions))) / len(actions) #start
208 equiprobable
209 for i in range(5000):
210     episode = generate_episode(policy2)
211     G = 0

```

```

207     visited = set()
208     for s,a,r in reversed(episode):
209         G = gamma*G + r
210         if (s,a) not in visited:
211             returns2[s,a] += 1
212             Q2[s,a] += (G - Q2[s,a]) / returns2[s,a]
213             best_a = np.argmax(Q2[s])
214             # $\epsilon$ -soft policy update: mostly greedy, small prob for others
215             for a_idx in range(len(actions)):
216                 if a_idx == best_a:
217                     policy2[s,a_idx] = 1 - epsilon + epsilon/len(actions)
218                 else:
219                     policy2[s,a_idx] = epsilon/len(actions)
220             visited.add((s,a))
221
222     print("\nMC with  $\epsilon$ -soft policy - learned policy:")
223     print(np.argmax(policy2, axis=1).reshape((grid_size, grid_size)))
224     plot_policy(np.argmax(policy2, axis=1), "MC  $\epsilon$ -soft")
225     V_mc_soft = np.max(Q2, axis=1)
226     print("\nValue function from MC with  $\epsilon$ -soft:")
227     print(np.round(V_mc_soft.reshape((grid_size, grid_size)),2))
228
229     #(3) Off-policy MC with importance sampling
230     target_policy = np.zeros((n_states, len(actions))) #target policy we want
231     #to learn: greedy in Q
232     Q3 = np.zeros((n_states, len(actions))) #action-value estimates
233     C = np.zeros((n_states, len(actions))) #cumulative sum of importance weights
234     #Initialize target policy as uniform
235     for s in range(n_states):
236         target_policy[s] = np.ones(len(actions)) / len(actions)
237     for i in range(5000):
238         episode = generate_episode(policy) #generate episode using behavior
239         #policy (here, policy is uniform)
240         G = 0
241         W = 1
242         for s,a,r in reversed(episode):
243             G = gamma*G + r
244             C[s,a] += W
245             Q3[s,a] += (W/C[s,a])*(G - Q3[s,a]) #weighted update
246             best_a = np.argmax(Q3[s])
247             #Make target policy greedy in Q
248             target_policy[s] = np.eye(len(actions))[best_a]
249             if a != best_a:
250                 break #stop if action taken is different from greedy:
251             #importance weight becomes 0
252             W = W / policy[s,a] #update importance weight
253     print("\nOff-policy MC with importance sampling - learned policy:")
254     print(np.argmax(target_policy, axis=1).reshape((grid_size, grid_size)))
255     plot_policy(np.argmax(target_policy, axis=1), "Off-policy MC")

```

Listing 2: Python Code for Monte Carlo with Exploring Starts,  $\epsilon$  - soft Policy, and Off-Policy Learning

## References

- [1] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.
- [2] Csaba Szepesvári. *Algorithms for Reinforcement Learning*. Morgan & Claypool Publishers, 2010.